

```
# Import relevant packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

# just to get rid of all those warning messages
import warnings
warnings.filterwarnings('ignore')

# Load the dataset
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive
```

Goal

The goal of this exploration is to create an optimal classifier that predicts if inputted Instagram accounts are fake or real. This classifier is built using Logistic Regression Models. I will make use of both a baseline Linear Regression Model and a hyperparameter tuned linear regression model to classify unseen accounts as real or fake.

Data Exploration

Link to Dataset: <https://www.kaggle.com/free4ever1/instagram-fake-spammer-genuine-accounts?select=test.csv>

```
train_df = pd.read_csv('/content/drive/MyDrive/Colab
Notebooks/ig_train.csv')
train_df.head()
```

| | profile pic | nums/length | username | ... | #follows | fake |
|---|-------------|-------------|----------|-----|----------|------|
| 0 | 1 | | 0.27 | ... | 955 | 0 |
| 1 | 1 | | 0.00 | ... | 533 | 0 |
| 2 | 1 | | 0.10 | ... | 98 | 0 |
| 3 | 1 | | 0.00 | ... | 651 | 0 |
| 4 | 1 | | 0.00 | ... | 126 | 0 |

```
[5 rows x 12 columns]
```

```
train_df.columns
```

```
Index(['profile pic', 'nums/length username', 'fullname words',
      'nums/length fullname', 'name==username', 'description length',
      'external URL', 'private', '#posts', '#followers', '#follows',
      'fake'],
      dtype='object')
```

```
# rename columns to omit whitespace
new_cols = {'profile pic': 'profile_pic', 'nums/length username':
'acc_name_ratio', 'fullname words': 'fullname_words',
'nums/length fullname': 'fullname_ratio',
'name==username': 'name_match', 'description length': 'bio_len',
'external URL': 'URL_flag', 'private': 'private_flag',
'#posts': 'post_nums', '#followers': 'follower_nums',
'#follows': 'follow_nums', 'fake': 'fake_flag'}
```

```
train_df.rename(columns=new_cols, inplace=True)
train_df.head()
```

| | profile_pic | acc_name_ratio | ... | follow_nums | fake_flag |
|---|-------------|----------------|-----|-------------|-----------|
| 0 | 1 | 0.27 | ... | 955 | 0 |
| 1 | 1 | 0.00 | ... | 533 | 0 |
| 2 | 1 | 0.10 | ... | 98 | 0 |
| 3 | 1 | 0.00 | ... | 651 | 0 |
| 4 | 1 | 0.00 | ... | 126 | 0 |

```
[5 rows x 12 columns]
```

```
# To understand the scale of our dataset
```

```
train_df.shape
```

```
(576, 12)
```

```
train_df.isna().sum()
```

```
profile_pic      0
acc_name_ratio   0
fullname_words   0
fullname_ratio   0
name_match       0
bio_len          0
URL_flag         0
private_flag     0
post_nums        0
follower_nums    0
follow:nums      0
fake_flag        0
dtype: int64
```

We have no null values in this dataset. Cleaning will be a lot simpler.

```
train_df.fake_flag.value_counts()
```

```
1    288
0    288
Name: fake_flag, dtype: int64
```

The number of observations in each classification match too.

Exploration of Binary Features

find columns that are binary

```
binary_features = [feature for feature in train_df.columns if  
len(train_df[feature].unique()) == 2]  
binary_features
```

```
['profile_pic', 'name_match', 'URL_flag', 'private_flag', 'fake_flag']
```

Drop fake flag as binary feature

```
binary_features = binary_features[:-1]
```

sum of binary variables/number of observations in class for real and fake accounts

```
bin_dist = train_df.groupby(by=train_df.fake_flag)  
[binary_features].sum()/train_df.groupby(by=train_df.fake_flag)  
[binary_features].count()  
bin_dist
```


| | profile_pic | name_match | URL_flag | private_flag |
|-----------|-------------|------------|----------|--------------|
| fake_flag | | | | |
| 0 | 0.993056 | 0.003472 | 0.232639 | 0.395833 |
| 1 | 0.409722 | 0.065972 | 0.000000 | 0.368056 |

visualization of ratio of binary features

```
ax = bin_dist.iloc[0].plot(kind='bar', color='blue', width=.75,  
legend=True, alpha=0.5)  
bin_dist.iloc[1].plot(kind='bar', color='orange', width=.5, alpha=0.7,  
legend=True)
```

```
plt.ylabel("Presence of Feature")  
plt.grid(which='major', alpha=0.7)  
plt.minorticks_on()  
plt.legend(["Real", "Fake"], loc='upper right')  
plt.xticks(rotation=45)
```

```
plt.title("Distribution of Binary Features Across Fake and Real  
Profiles")  
plt.show()
```



Surprisingly, it looks like the presence of a URL is indicative of a real account

Exploration of Other Features


Look at description length across real and fake profiles

```
desc_dist_real = train_df.loc[train_df.fake_flag == 0].bio_len.copy()
desc_dist_fake = train_df.loc[train_df.fake_flag == 1].bio_len.copy()
```

```
ax = desc_dist_real.plot(kind='hist', color='blue', legend=True,
alpha=0.5)
desc_dist_fake.plot(kind='hist', color='orange', alpha=0.7,
legend=True)
```

```
plt.ylabel("Frequency")
plt.xlabel("Bio Length (characters)")
plt.grid(which='major', alpha=0.7)
plt.minorticks_on()
plt.legend(["Real", "Fake"], loc='upper right')
```

```
plt.title("Distribution of Bio Length Across Fake and Real Profiles")
plt.show()
```



Look at followers distribution across real and fake profiles

```
follower_dist_real = train_df.loc[train_df.fake_flag ==
0].follower_nums.copy()
follower_dist_fake = train_df.loc[train_df.fake_flag ==
1].follower_nums.copy()
```

Assessing the range for the two account types

```
print("Range for real account followers:
{:,.}").format(follower_dist_real.max() - follower_dist_real.min())
print("Range for fake account followers:
{:,.}").format(follower_dist_fake.max() - follower_dist_fake.min())
```

Range for real account followers: 15,338,529

Range for fake account followers: 3,033

As the difference in range for these columns is large, we standardize the data

```
follower_dist_real_scaled = (follower_dist_real -
follower_dist_real.mean())/(follower_dist_real.std())
follower_dist_fake_scaled = (follower_dist_fake -
follower_dist_fake.mean())/(follower_dist_fake.std())
```

```
ax = follower_dist_real_scaled.plot(kind='hist', color='blue',
legend=True, alpha=0.5)
follower_dist_fake_scaled.plot(kind='hist', color='orange', alpha=0.7,
legend=True)
```

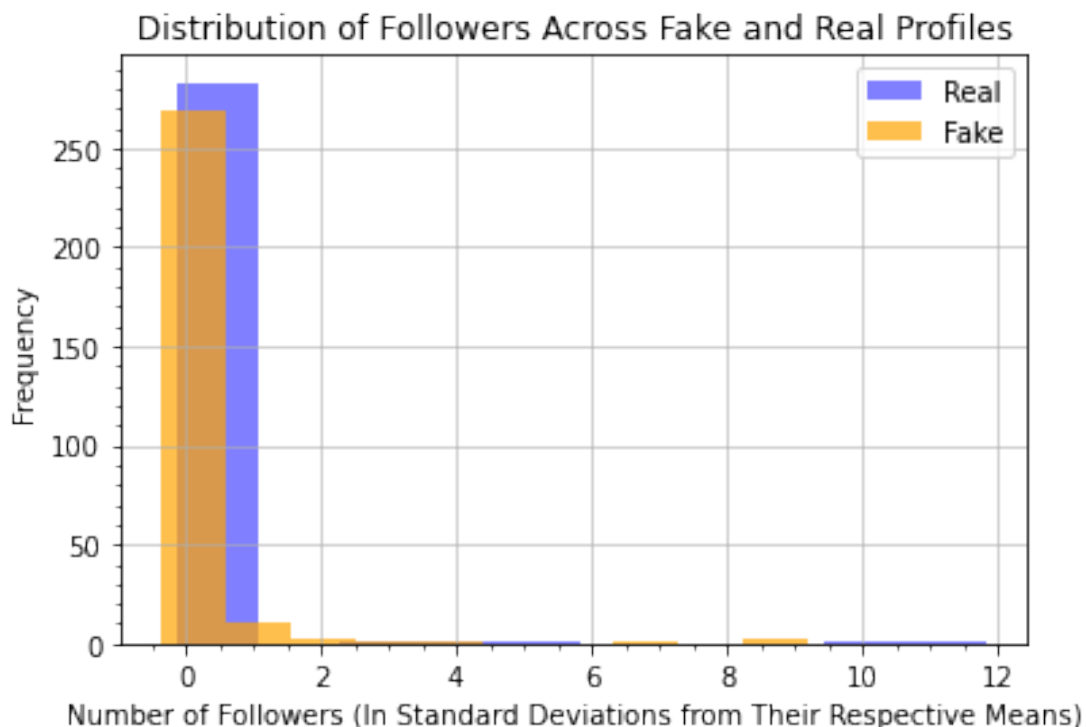
```
plt.ylabel("Frequency")
```

```

plt.xlabel("Number of Followers (In Standard Deviations from Their  
Respective Means)")
plt.grid(which='major', alpha=0.7)
plt.minorticks_on()
plt.legend(["Real", "Fake"], loc='upper right')

plt.title("Distribution of Followers Across Fake and Real Profiles")
plt.show()
print("Mean Followers of Real Accounts:  
{:,.0F}".format(follower_dist_real.mean()))
print("Mean Followers of Fake Accounts:  
{:,.0F}".format(follower_dist_fake.mean()))

```



Mean Followers of Real Accounts: 170,504
Mean Followers of Fake Accounts: 111

Mean followers for real accounts seems high. We must consider the data collection process of this dataset as it appears that the dataset may not be represent the population well.

```

# Look at following distribution across real and fake profiles
follow_dist_real = train_df.loc[train_df.fake_flag ==
0].follow_nums.copy()
follow_dist_fake = train_df.loc[train_df.fake_flag ==
1].follow_nums.copy()

# Assessing the range for the two account types
print("Range for real account follows:  
{:,}".format(follow_dist_real.max() - follow_dist_real.min()))

```

```

print("Range for fake account follows:
{:,.}").format(follow_dist_fake.max() - follow_dist_fake.min())

Range for real account follows: 7,500
Range for fake account follows: 7,272


# As the range for total follows is similar to each other, we don't
need to scale

ax = follow_dist_real.plot(kind='hist', color='blue', legend=True,
alpha=0.5)
follow_dist_fake.plot(kind='hist', color='orange', alpha=0.7,
legend=True)

plt.ylabel("Frequency")
plt.xlabel("Number of Follows")
plt.grid(which='major', alpha=0.7)
plt.minorticks_on()
plt.legend(["Real", "Fake"], loc='upper right')

plt.title("Distribution of Follows Across Fake and Real Profiles")
plt.show()
print("Mean Follows of Real Accounts:
{:,.0F}").format(follow_dist_real.mean())
print("Mean Follows of Fake Accounts:
{:,.0F}").format(follow_dist_fake.mean())

```



Mean Follows of Real Accounts: 715
Mean Follows of Fake Accounts: 302

```
# Look at following distribution across real and fake profiles
post_dist_real = train_df.loc[train_df.fake_flag ==
0].post_nums.copy()
post_dist_fake = train_df.loc[train_df.fake_flag ==
1].post_nums.copy()

# Assessing the range for the two account types
print("Range for real account posts: {:,}".format(post_dist_real.max()
- post_dist_real.min()))
print("Range for fake account posts: {:,}".format(post_dist_fake.max()
- post_dist_fake.min()))
```


Range for real account posts: 7,389
Range for fake account posts: 299

```
# As the difference in range for these columns is large, we
standardize the data
post_dist_real_scaled = (post_dist_real -
post_dist_real.mean())/(post_dist_real.std())
post_dist_fake_scaled = (post_dist_fake -
post_dist_fake.mean())/(post_dist_fake.std())

ax = post_dist_real_scaled.plot(kind='hist', color='blue',
legend=True, alpha=0.5)
post_dist_fake_scaled.plot(kind='hist', color='orange', alpha=0.7,
legend=True)
```

```
plt.ylabel("Frequency")
plt.xlabel("Number of Posts (In Standard Deviations from Their
Respective Means)")
plt.grid(which='major', alpha=0.7)
plt.minorticks_on()
plt.legend(["Real", "Fake"], loc='upper right')

plt.title("Distribution of Posts Across Fake and Real Profiles")
plt.show()
print("Mean Posts of Real Accounts:
{:,.0F}".format(post_dist_real.mean()))
print("Mean Posts of Fake Accounts:
{:,.0F}".format(post_dist_fake.mean()))
```

Mean Posts of Real Accounts: 206

Mean Posts of Fake Accounts: 9

Similar to number of followers, the mean posts also seems high.

Implemetation of Logistic Regression to Determine Profile Type

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, plot_confusion_matrix
```

As all datapoints are numeric, no data encoding is required

Split output and training data

```
train_outputs = train_df.fake_flag
train_df.pop('fake_flag')
```

```
0      0
1      0
2      0
3      0
4      0
...
571    1
572    1
573    1
574    1
```

575 1

Name: fake_flag, Length: 576, dtype: int64

```
def evaluate(model, X, y_true):
    '''
    Calculates the accuracy for a trained logistic regression model

    Parameters:
    model: Trained Linear Regression Model
    X, pandas.DataFrame -- test/validation input
    y_true, pandas.DataFrame, numpy.ndarray -- real output for X

    Returns:
    accuracy, float -- accuracy of the model on passed data
    auc, float -- Area under the ROC (score for model performance)
    '''

    # Calculate predicted values
    y_pred = model.predict(X)

    correct = 0

    # if predicted val matches real val, increase correct
    for i, pred in enumerate(y_pred):
        if pred == y_true.iloc[i]:
            correct += 1

    # Calculate accuracy
    accuracy = (correct / y_true.shape[0]) * 100

    # Calculate AUC
    auc = roc_auc_score(y_true, y_pred)

    return accuracy, auc
```

Note: In this situation, the accuracy and the AUC are roughly the same. This is because the number of observation in each class for the training set is the same.

Implementation of a Baseline Logistic Regression Classifier

Initialise test data

```
test_df = pd.read_csv('/content/drive/MyDrive/Colab  
Notebooks/ig_test.csv')
```

```
test_df.rename(columns=new_cols, inplace=True)
```

```
test_outputs = test_df.fake_flag
```

```
test_df.pop('fake_flag')
```

```
0     0
1     0
2     0
```

```
3      0
4      0
      ..
115    1
116    1
117    1
118    1
119    1
```

```
Name: fake_flag, Length: 120, dtype: int64
```

```
# Initialize and Fit Model
```

```
model = LogisticRegression()
```

```
model.fit(train_df, train_outputs)
```

```
# Obtain Evaluation Metrics
```

```
acc, auc = evaluate(model, test_df, test_outputs)
```

```
print('Model AUC: {:.2f}'.format(auc))
```

```
Model AUC: 0.908333
```

```
y_pred = model.predict(test_df)
```

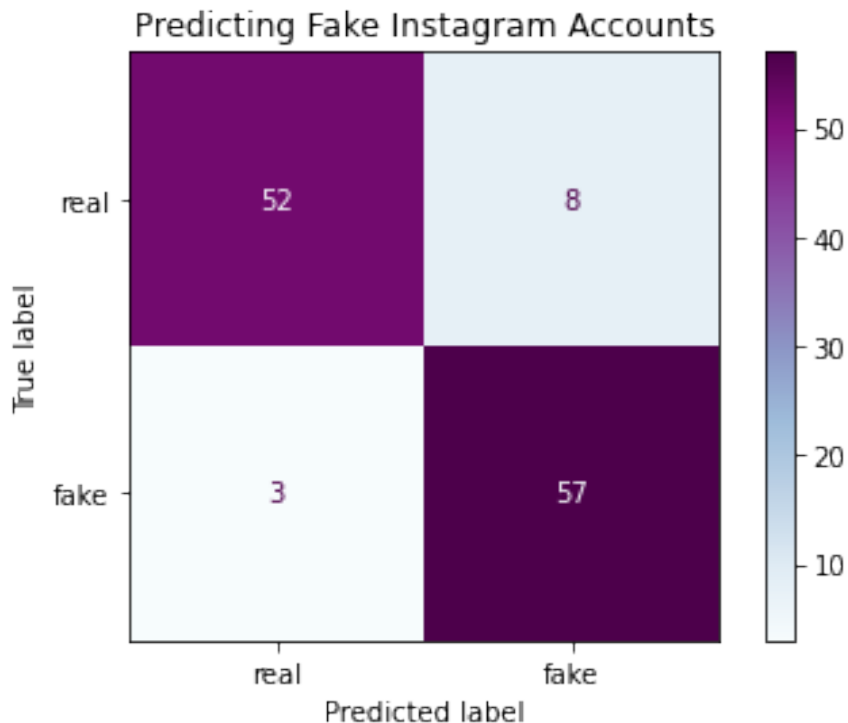
```
# Plotting confusion matrix for model outputs
```

```
labels = ["real", "fake"]
```

```
plot_confusion_matrix(model, test_df, test_outputs,  
display_labels=labels, cmap='BuPu')
```

```
plt.title("Predicting Fake Instagram Accounts")
```

```
Text(0.5, 1.0, 'Predicting Fake Instagram Accounts')
```



Feature Importance for Baseline Model

feature weights

```
weights = model.coef_[0]
```

Dataframe storing feature name and it's assigned weight

```
feat_importance = pd.DataFrame([train_df.columns, weights]).T
```

Rename Columns

```
feat_importance.rename(columns={0:'Feature', 1:'Weight'},  
inplace=True)
```

Sort from highest weight to lowest and display

```
feat_importance = feat_importance.sort_values(by='Weight',  
ascending=False, key=abs)
```

```
feat_importance.reset_index(drop=True, inplace=True)
```

```
display(feat_importance)
```

| | Feature | Weight |
|---|----------------|------------|
| 0 | acc_name_ratio | 3.34628 |
| 1 | profile_pic | -2.95199 |
| 2 | private_flag | -0.976308 |
| 3 | fullname_ratio | 0.806112 |
| 4 | name_match | 0.75261 |
| 5 | URL_flag | -0.746025 |
| 6 | fullname_words | -0.478901 |
| 7 | post_nums | -0.0140987 |
| 8 | bio_len | -0.0120212 |

```
9     follower_nums    -0.0026903
10     follow_nums     0.000789614
```

Logistic Regression Classifier After Tuning

A Validation set is needed because we want to tune our LRM's hyperparameters depending on the AUC on the validation set

```
x_train, x_valid, y_train, y_valid = train_test_split(train_df,
train_outputs, test_size=0.2, random_state=0)
```

The model is tuned using penalty_types, and C_values

```
def tune_logistic_regression(X_train, y_train, X_valid, y_valid,
penalty_types, C_values, verbose=True):
```

```
    '''Tunes logistic regression models over the hyperparameters
penalty type and C
to maximize the accuracy of the models
```

```
    Parameters:
```

```
    X_train: pandas.DataFrame -- Input for training data
```

```
    y_train: pandas.DataFrame, numpy.ndarray -- Output for training
```

```
data
```

```
    X_valid: pandas.DataFrame -- Input for validation data
```

```
    y_valid: pandas.DataFrame, numpy.ndarray -- Output for
```

```
validation data
```

```
    penalty_types, str -- norm tpye ('l1', 'l2')
```

```
    C_values, numpy.ndarray, list(int) -- Regularization Strength
```

```
    Returns:
```

```
    best_model: sklearn.LinearRegression -- Trained model with the
best AUC
```

```
    parameters: dict -- dictionary of the best parameters
```

```
    val_aucs, numpy.ndarray -- array holding the AUC of every model
```

```
    '''
```

Create empty lists where we will store the results of hyperparameter tuning

```
parameters = []
```

```
models = []
```

```
val_aucs = []
```

Loop through the hyperparameters of interest

```
for penalty in penalty_types:
```

```
    for C in C_values:
```

Train the logistic regression model with the given hyperparameters

```
        lr = LogisticRegression(C=C, penalty=penalty,
solver='liblinear')
```

Fit the model using the training data

```

lr = lr.fit(X_train, y_train)

# Get the evaluation metrics on the validation set
acc, auc = evaluate(lr, X_valid, y_valid)

# Store the results
parameters.append({'penalty': penalty, 'C': C})
models.append(lr)
val_aucs.append(auc)

print('--'*20, '\nModel Parameters:\nPenalty: {}, C: {}'.format(penalty, C))
print('Model AUC: {:.2f}'.format(auc))

# Determine the best model -- that is, the one with the greatest AUC
best_model_index = np.argmax(val_aucs)
best_model = models[best_model_index]

print('- - '*10, "\nBest model parameters: ",
parameters[best_model_index])
print("Best model AUC: ", val_aucs[best_model_index])

# Return best model
return best_model, parameters, val_aucs

penalty_types=["l1", "l2"]
C_values=[0.01, 0.1, 1, 10, 100, 1000]

final_model, params_list, aucs = tune_logistic_regression(x_train,
y_train, x_valid, y_valid, penalty_types, C_values, verbose=True)

-----
Model Parameters:
Penalty: l1, C: 0.01
Model AUC: 0.724138
-----
Model Parameters:
Penalty: l1, C: 0.1
Model AUC: 0.870690
-----
Model Parameters:
Penalty: l1, C: 1
Model AUC: 0.922414
-----
Model Parameters:
Penalty: l1, C: 10
Model AUC: 0.913793
-----

```

```

Model Parameters:
Penalty: l1, C: 100
Model AUC: 0.913793
-----
Model Parameters:
Penalty: l1, C: 1000
Model AUC: 0.913793
-----
Model Parameters:
Penalty: l2, C: 0.01
Model AUC: 0.732759
-----
Model Parameters:
Penalty: l2, C: 0.1
Model AUC: 0.732759
-----
Model Parameters:
Penalty: l2, C: 1
Model AUC: 0.732759
-----
Model Parameters:
Penalty: l2, C: 10
Model AUC: 0.732759
-----
Model Parameters:
Penalty: l2, C: 100
Model AUC: 0.732759
-----
Model Parameters:
Penalty: l2, C: 1000
Model AUC: 0.732759
-----
Best model parameters: {'penalty': 'l1', 'C': 1}
Best model AUC: 0.9224137931034483

```

It is important to note that the performance above is only on the validation set. The reason we measure AUC for the validation set is to see what parameters work best for this situation prior to exposing our model to test data. Now that we have the ideal model and parameters, we can assess the performance of our model on the test data

```

acc, auc = evaluate(final_model, test_df, test_outputs)

print('Model AUC: {:.2f}'.format(auc))

Model AUC: 0.891667

y_pred = final_model.predict(test_df)

# Plotting confusion matrix for model outputs
labels = ["real", "fake"]


```

```

plot_confusion_matrix(final_model, test_df, test_outputs,
display_labels=labels, cmap='BuPu')
plt.title("Predicting Fake Instagram Accounts")

Text(0.5, 1.0, 'Predicting Fake Instagram Accounts')

```



Feature Importance for Best Tuned Model

```

# feature weights
weights = final_model.coef_[0]

# Dataframe storing feature name and it's assigned weight
feat_importance2 = pd.DataFrame([train_df.columns, weights]).T
# Rename Columns
feat_importance2.rename(columns={0:'Feature', 1:'Weight'},
inplace=True)
# Sort from highest weight to lowest and display
feat_importance2 = feat_importance2.sort_values(by='Weight',
ascending=False, key=abs)
feat_importance2.reset_index(drop=True, inplace=True)

```

```
display(feat_importance2)
```

| | Feature | Weight |
|---|----------------|----------|
| 0 | acc_name_ratio | 6.32608 |
| 1 | profile_pic | -2.86889 |
| 2 | URL_flag | -1.71138 |
| 3 | name_match | 1.2862 |
| 4 | fullname_words | -0.24888 |


```
5     private_flag    -0.170155
6         post_nums   -0.0173395
7             bio_len -0.00745825
8     follower_nums   -0.00207805
9         follow_nums  0.00100328
10  fullname_ratio      0
```

Evaluation

While the best tuned model works very well on the validation set, the baseline model performs better overall on the test set. This is interesting because one would expect that models that have been tuned would procure the most accurate predictions.

One possible reason for this is that we do not have an extensive training sample. As the training set consists of just 576 observations, splitting that down into training and validation could result in an underfitting of training data for the models that are tuned. Therefore, these models are not as generalizable.

Possible Extensions to the Project

I might try bootstrapping training data to get a larger training sample and make the output more generalizable. I believe that a Bootstrapped Aggeragate could help combat this issue of underfitting and provide a more robust classifier